

Introduction To Database

CT042-3-1-IDB

0004 - Data Model

Topic & Structure of The Lesson

- Data Model
- Business Rules
- Relational Model
- Entity Relationship Model

Learning Outcomes

At the end of this topic, You should be able to

- Describe types of data model
- Write business rules
- Understand entity relationship model

Key Terms You Must Be Able To Use

If you have mastered this topic, **you should be able to use the following terms correctly in your assignments and exams:**

- Data model
- Business rules
- ERD

What is a Data Model?

1. Definition

A data model is a **framework** or a **blueprint** that **visually represents** and **organizes data**, **illustrating how different data elements interact within a system**. It helps in understanding and structuring complex data, enabling efficient storage, retrieval, and management of information.

A data model is essentially a simplified representation of real-world data structures. It **uses graphical symbols and relationships to depict data elements and their connections**, making it easier to understand and manage complex data environments.

Data models facilitate communication among various stakeholders, including:

- **Designers** who create the data structure.
- **Applications Programmers** who implement and interact with the data.
- **End Users** who utilize the data for practical purposes.

2. Key Characteristics of Data Models

Relatively Simple Representations:

- **Graphical Depictions:** Data models are often visualized through diagrams that simplify complex data structures. For example, Entity-Relationship Diagrams (ERDs) use shapes like rectangles for entities and diamonds for relationships to illustrate how different data elements interact.
- **Abstraction:** They abstract and simplify real-world data complexities into a manageable form. For instance, a data model for a library system might represent books, authors, and borrowers as distinct entities and show their relationships.

Facilitating Interaction:

- **Designer:** Helps designers understand how to structure the data in a way that aligns with business needs.
- **Applications Programmer:** Provides a blueprint for programmers to implement data storage and retrieval processes efficiently.
- **End User:** Ensures that the data representation aligns with user requirements and makes data access intuitive.

2. Key Characteristics of Data Models

Visual Representation of Data Connections:

- **Structure and Relationships:** Shows how different data elements relate to each other. For instance, in a customer relationship management (CRM) system, a data model might show how customer information is connected to sales orders and support tickets.

The Importance of Data Models

1. Facilitating Communication

Data models act as a communication tool among various stakeholders involved in database development. They help bridge the gap between different perspectives, ensuring that everyone has a unified understanding of how data should be structured and used.

Example:

- Imagine a company is developing a new inventory management system. The database designer creates a data model that outlines how inventory data will be stored and related to other data like suppliers and sales.
- The applications programmer uses this model to build the system's interfaces and reports. Meanwhile, the inventory manager and purchasing manager refer to this model to understand how their respective data concerns are addressed within the system.
- Without a clear data model, these stakeholders might have conflicting views on data management, leading to inefficiencies and errors.

2. Improving Organizational Understanding

A well-developed data model can enhance the understanding of an organization's data by mapping out how various elements of the business are interconnected. This is particularly valuable for stakeholders who need to grasp the bigger picture of how different parts of the organization fit together.

Example:

- A business owner who has been deeply involved in day-to-day operations might have a detailed understanding of their specific department but lack a comprehensive view of the entire organization.
- By reviewing a well-structured data model, they can see how data flows between departments, such as how sales data impacts inventory and how inventory levels affect purchasing decisions. This holistic view can lead to better strategic decisions.

3. Different Perspectives on Data

Different users and producers of data have varying needs and perspectives on how data should be managed and used. Data models help accommodate these diverse viewpoints by providing a common framework for data representation.

Example:

- The company president requires a high-level view of data to make strategic decisions, such as overall sales performance and financial health.
- On the other hand, a purchasing manager needs detailed information about suppliers and item costs. A data model accommodates these differing needs by defining various data structures and relationships that support both high-level and detailed analyses.

4. Preventing Conflicts and Errors

A good data model helps prevent conflicts and inconsistencies in how data is managed and interpreted. Without a coherent data model, different systems or departments might use incompatible data formats or definitions, leading to errors and inefficiencies.

Example:

- Suppose an inventory management system and an order entry system use different product-numbering schemes. If these systems are not aligned, it could result in order errors, incorrect inventory levels, and financial discrepancies.
- A data model helps ensure that all systems use consistent data definitions and formats, reducing the risk of such problems.

5. Blueprint for System Development

Just as a house blueprint provides a plan for construction, a data model provides a plan for database development. It outlines the structure and relationships of data, guiding the development of databases and applications.

Example:

- Before building a new customer relationship management (CRM) system, a company creates a data model that defines how customer information, interactions, and transactions are organized.
- This model serves as a guide for developers to create a system that effectively manages customer data and supports business processes.
- Without this blueprint, the development process could lack direction, resulting in a disjointed and ineffective system.

Conclusion

- Data models are essential for effective database design and management. They facilitate communication among stakeholders, enhance organizational understanding, accommodate diverse perspectives, prevent conflicts, and serve as a blueprint for system development.
- By providing a clear and structured view of data, data models help ensure that databases and applications are well-organized and aligned with business needs.

Data Model Basic Building Blocks

1. Definition

The term Data Model Basic Building Blocks refers to the fundamental components that constitute a data model.

These components are essential for structuring, organizing, and managing data within a database system.

They help define how data is represented, related, and constrained.

The basic building blocks of data models include entities, attributes, relationships, and constraints.

2. Entities

2A. Definition:

- **An entity is essentially an object or concept about which you want to collect and store information. It can be a physical item, a person, a concept, or even an event.**
- Entities are the fundamental objects in a data model that have attributes (properties or details) and can have relationships with other entities.

2B. Characteristics of Entities:

- **Distinct:** Each entity represents a unique object or concept with a distinct identity.
- **Has Attributes:** Entities are characterized by attributes, which are the details or properties that describe them.
- **Can Be Related:** Entities can have relationships with other entities, indicating how they interact or are associated with one another.

2. Entities

2C. Examples of Entities

Entity: Employee

Description: Represents individuals who work for a company.

Attributes:

- **EmployeeID:** A unique identifier for the employee.
- **Name:** The employee's full name.
- **Position:** Job title of the employee.
- **Department:** The department in which the employee works.
- **Salary:** The employee's salary.

Use Case: In a company's human resources system, you would collect and store data about each employee, such as their contact information, job history, and performance reviews.

2. Entities

2D. Examples of Entities

Entity: Product

Description: Represents items that a company sells.

Attributes:

- **ProductID:** A unique identifier for the product.
- **Name:** The name of the product.
- **Price:** The selling price of the product.
- **Category:** The category or type of product (e.g., electronics, clothing).
- **StockQuantity:** The number of units available in inventory.

Use Case: In an e-commerce database, each product would be an entity, with attributes detailing its characteristics and availability.

2. Entities

2E. Examples of Entities

Entity: Customer

Description: Represents individuals or organizations that purchase goods or services.

Attributes:

- **CustomerID:** A unique identifier for the customer.
- **Name:** The name of the customer.
- **Email:** The customer's email address.
- **PhoneNumber:** Contact number of the customer.
- **Address:** Shipping address for deliveries.

Use Case: In a CRM (Customer Relationship Management) system, you track customer interactions, orders, and preferences using customer entities.

2. Entities

2F. Examples of Entities

Entity: Course

Description: Represents a specific educational course offered by an institution.

Attributes:

- **CourseID:** A unique identifier for the course.
- **Title:** The name or title of the course.
- **Credits:** Number of credits awarded upon completion.
- **Instructor:** The person teaching the course.
- **Schedule:** The time and days when the course is offered.

Use Case: In a university's course management system, each course would be an entity, with attributes describing the course content, schedule, and associated faculty.

3. Attributes

3A. Definition:

- **Attributes are characteristics or properties that describe an entity.** They provide details about each entity.

3B. Examples:

- **Customer Attributes:** Last Name, First Name, Phone Number, Address, and Credit Limit.
- **Product Attributes:** Product Name, Price, Category, and Stock Quantity.
- **Event Attributes:** Event Name, Date, Location, and Organizer.

3C. Explanation:

- **Attributes** are analogous to **fields** or **columns in a table**. They define the specifics of each entity, allowing detailed data storage and retrieval.

4. Relationship

4A. Definition:

- Relationships describe **how entities are associated with each other**. They illustrate how data in one entity is related to data in another.
- Relationships help define how entities interact and connect within the data model, ensuring that the structure reflects real-world associations.

4B. Types and Examples:

One-to-Many (1:M or 1..*):

- A single instance of an entity is associated with multiple instances of another entity.
- **Example:** A Painter (one) can create many Paintings (many). This relationship is denoted as "PAINTER paints PAINTING."
- **Example:** A Customer (one) can generate many Invoices (many). This is labeled "CUSTOMER generates INVOICE."

4. Relationship

4B. Types and Examples:

Many-to-Many (M:N or *..*):

- Multiple instances of one entity are associated with multiple instances of another entity.
- **Example:** An Employee (many) can learn many Skills (many), and each Skill (many) can be learned by many Employees. This relationship is denoted as "EMPLOYEE learns SKILL."
- **Example:** A Student (many) can enroll in many Classes (many), and each Class (many) can have many Students. This is labeled "STUDENT takes CLASS."

4. Relationship

4B. Types and Examples:

One-to-One (1:1 or 1..1):

- A single instance of one entity is associated with a single instance of another entity.
- **Example:** Each Store (one) is managed by one Employee (one), and each Employee (one) manages only one Store. This relationship is denoted as "EMPLOYEE manages STORE."

5. Constraints

5A. Definition:

- Constraints are **rules or restrictions/limitations applied to data to maintain its integrity and validity**.

5B. Examples:

- **Salary Constraint:** An employee's salary must be between \$6,000 and \$350,000. This ensures that data is within a realistic and acceptable range.
- **GPA Constraint:** A student's GPA must be between 0.00 and 4.00. This constraint maintains the integrity of academic performance data.
- **Class Constraint:** Each class must have exactly one teacher. This ensures that every class is assigned to a single instructor.

5C. Explanation:

- Constraints are crucial for maintaining data quality and ensuring that the data adheres to predefined rules. They prevent invalid or inconsistent data entries.

6. Identifying Entities, Attributes, Relationships, and Constraints

6A. Steps:

- **Understand Business Rules:** Begin by thoroughly understanding the business processes and requirements of the domain you are modeling. This involves gathering requirements and identifying the key elements that need to be represented.
- **Define Entities:** Identify the distinct objects or concepts that need to be tracked. Determine what each entity represents and how it will be uniquely identified.
- **Specify Attributes:** List the characteristics or properties that need to be recorded for each entity.
- **Establish Relationships:** Determine how entities interact with each other and define the nature of these interactions.
- **Apply Constraints:** Define rules to ensure data integrity and consistency, reflecting real-world restrictions.

By following these steps, you can create a robust and effective data model that accurately represents the data and relationships within a specific domain.

Business Rules in Database Design

1. What are Business Rules?

Business rules **are precise descriptions of policies, procedures, or principles within an organization.**

They define how business operations should be conducted and play a crucial role in ensuring consistency and clarity across all processes.

Apply to any organization that stores and uses data to generate information.

In the context of database design, business rules are essential because they guide the structure and constraints of the data model.

Business rules are the foundation for creating meaningful and consistent databases that align with an organization's operations. They define how data should be structured, what constraints should be applied, and how different entities within the organization relate to each other. By adhering to these rules, organizations ensure that their data is accurate, reliable, and useful for decision-making.

2. Example of Business Rules

Pilot Duty Rule: A pilot cannot be on duty for more than 10 hours during a 24-hour period.

- **Entity:** Pilot
- **Attribute:** Duty Hours
- **Constraint:** Maximum of 10 hours within 24 hours
- **Relationship:** A pilot's duty is constrained by time, ensuring safety and compliance with regulations.

Professor Teaching Rule: A professor may teach up to four classes during a semester.

- **Entity:** Professor, Class
- **Attribute:** Number of Classes
- **Constraint:** Maximum of 4 classes per semester
- **Relationship:** A professor is linked to the classes they teach, and this relationship is limited by the rule.

3. Importance in Database Design

When database designers are building a data model, they need to consider the types of data within the organization, how the data is used, and the timeframes involved.

However, understanding the organization's operations through business rules is critical for the data to be meaningful. Business rules help in defining entities (such as a Pilot or Professor), attributes (like Duty Hours or Number of Classes), and relationships (such as a Professor teaching Classes).

Without business rules, the data could become inconsistent or not align with the organization's needs.

For example, if a business rule about the maximum number of classes a professor can teach is not enforced, a professor might be assigned too many classes, leading to scheduling conflicts and reduced teaching quality.

4. How Business Rules Are Applied

Business rules are used to enforce actions within an organization's environment. For example, the rule about pilot duty hours ensures that pilots are not overworked, which is crucial for safety.

Similarly, limiting the number of classes a professor can teach ensures they can manage their workload effectively.

In a database, these rules translate into constraints and relationships:

- **Constraints:** Ensure that the data adheres to the rules (e.g., the maximum number of hours a pilot can be on duty).
- **Relationships:** Define how entities interact (e.g., a professor can teach multiple classes, but each class is linked to only one professor).

5. Writing and Updating Business Rules

Business rules must be documented clearly and updated regularly to reflect any changes in the organization's operations.

For instance, if a new regulation requires that pilots must rest for 12 hours before their next duty, the business rule and corresponding database constraint must be updated to enforce this.

6. Practical Example in a Database

Let's consider a business rule: **A customer may generate many invoices, but an invoice is generated by only one customer.**

- **Entities:** CUSTOMER and INVOICE
- **Relationship:** 1:M (One-to-Many) where one customer can generate multiple invoices.
- **Constraint:** Each invoice is associated with exactly one customer.

In the database, this rule would be reflected by having a Customer ID in the Invoice table as a foreign key. The database would enforce that every invoice must have a valid customer ID, ensuring that invoices are always linked to a specific customer.

7. Where Can You Get The Business Rules?

Business rules are crucial in database design because they define the guidelines and constraints that govern how data is collected, stored, and processed within an organization. They come from various sources, each contributing different perspectives and insights.

7A. Sources of Business Rules

i. **Company Managers and Policy Makers:**

Example: A company manager might establish a rule that no purchase order can be processed without a manager's approval if it exceeds a certain amount. This rule ensures financial control and accountability within the organization.

Why Important: Managers and policy makers set the strategic direction and high-level guidelines that shape the company's operations. Their rules are often tied to compliance, legal requirements, and strategic objectives.

7. Where Can You Get The Business Rules?

7A. Sources of Business Rules

ii. Department Managers:

Example: A sales department manager might define a rule that customer orders should be processed within 24 hours of receipt to ensure timely delivery. This rule is important for maintaining customer satisfaction.

Why Important: Department managers provide rules that are specific to their functional areas. These rules help ensure that departmental processes align with overall company goals and operate efficiently.

iii. Written Documentation:

Example: A company's operations manual might detail the steps for handling customer complaints, including the requirement that all complaints be logged within a database within 2 hours of receipt.

Why Important: Written documentation such as **procedures, standards, and operations manuals** provide a consistent and formalized source of business rules. They are essential for ensuring that all employees follow the same procedures and for preserving institutional knowledge.

iv. Interviews with End Users:

Example: Direct interviews with end users, such as customer service representatives, might reveal that they prioritize responding to high-priority customer tickets first. This could lead to a rule that flags high-priority tickets in the system for immediate attention.

Why Important: End users interact directly with the systems and processes on a daily basis. They offer practical insights into how business rules are applied in real-world scenarios, which might not be captured in formal documentation.

7. Where Can You Get The Business Rules?

7B. Challenges with End-User Perceptions

While end users are invaluable for providing insights, their perceptions can sometimes be misleading or inconsistent:

Example: A situation might arise where a warehouse worker believes that they can override the inventory management system to adjust stock levels, but company policy might state that only supervisors have this authority. If the database is designed based on the worker's incorrect assumption, it could lead to unauthorized changes in inventory data, affecting the accuracy of stock levels and potentially leading to stockouts or overstocking.

Reconciliation and Verification: To avoid such issues, it's crucial for the database designer to reconcile different perceptions by cross-referencing the information from multiple sources, including policy documents and interviews with managers. This process ensures that the business rules captured are accurate and consistent with company policies.

7C. Importance of Documenting Business Rules

Documenting business rules is not just a formality; it serves multiple purposes in database design:

Standardization of Data View:

- **Example:** If different departments have varied definitions of what constitutes a "customer," the database could end up with inconsistent data. Standardizing the definition through business rules ensures that everyone in the company uses the term in the same way, leading to more reliable data.

Communication Tool:

- **Example:** A documented rule stating that "all invoices must be archived for at least 7 years" communicates clear expectations to both users and designers, ensuring that the database is designed to store this data accordingly.

Understanding of Data Nature and Scope:

- **Example:** By knowing that a business rule requires tracking customer order history for 10 years, a designer can better understand the volume and type of data that will need to be managed.

7. Where Can You Get The Business Rules?

7C. Importance of Documenting Business Rules

Understanding Business Processes:

- **Example:** If a business rule dictates that all returns must be processed within 5 days, it indicates that the database must support tracking and managing the return process efficiently.

Development of Relationship Participation Rules:

- **Example:** A rule might specify that each sales transaction must be linked to exactly one customer account. This relationship is crucial for maintaining data integrity and ensuring that sales data is properly attributed.

7D. Limitations of Modeling Business Rules

Not all business rules can be directly modeled within the database. Some rules, particularly those involving temporal or external factors, must be enforced by the application layer:

- **Example:** The rule "no pilot can fly more than 10 hours within any 24-hour period" involves tracking time and potentially data from external systems (like flight logs). While the database can store the necessary data (flight times), the logic to enforce this rule must be implemented in the application software that interacts with the database.

Conclusion

Understanding where business rules come from and how to accurately document and implement them is fundamental to successful database design. It ensures that the database not only meets the technical requirements but also aligns with the business's operational needs and constraints. By addressing potential discrepancies and limitations, the designer can create a system that is both robust and flexible, capable of supporting the organization's goals effectively.

8. Translating Business Rules into Data Model Components

When translating business rules into data model components, the process begins by identifying the key elements within the business environment. Business rules help define how these elements—such as people, objects, or events—interact. This interaction is critical because it shapes the data model's structure, ensuring that it accurately reflects the real-world system it represents.

8A. Identifying Entities and Relationships

Entities are the fundamental objects of interest in the data model. Typically, these are represented by **nouns** in business rules. For example, in the rule "a customer may generate many invoices," the nouns "customer" and "invoices" represent entities because they are objects the business environment needs to track.

Relationships represent the way these entities interact with each other, often indicated by **verbs** in the business rules. In the example, "generate" is the verb that defines the relationship between the "customer" and "invoices."

8B. Example: Customer and Invoices

Consider the business rule:

- "A customer may generate many invoices."

Here, you identify two entities:

- **Customer** (an individual or company that purchases goods or services).
- **Invoice** (a document that records a transaction).

The **verb** "generate" indicates a relationship between these entities.

In this case, the relationship is one-to-many (1:M) because:

- **One customer** can generate **many invoices** (Customer to Invoice relationship).
- **One invoice** is generated by **only one customer** (Invoice to Customer relationship).

8. Translating Business Rules into Data Model Components

8C. Determining the Type of Relationship

Relationships in a data model are bidirectional (A Entity to B Entity; B Entity to A Entity), meaning they can be viewed from both perspectives.

To accurately identify the type of relationship, you should ask:

- **How many instances of Entity B (Invoices) are related to one instance of Entity A (Customer)?**
 - **Answer:** Many invoices can be generated by one customer.
- **How many instances of Entity A (Customer) are related to one instance of Entity B (Invoice)?**
 - **Answer:** One customer can generate many invoices.

The answers indicate a one-to-many (1:M) relationship.

8D. Example: Student and Class

Let's examine a different scenario:

- **Business rule:** "A student can enroll in many classes."

In this case:

- Student and Class are entities.
- The relationship involves the verb "enroll."

To determine the relationship type:

- **How many instances of Class are related to one instance of Student?**
 - **Answer:** A student can enroll in many classes (many classes to one student).
- **How many instances of Student are related to one instance of Class?**
 - **Answer:** Many students can enroll in one class (many students to one class).

This creates a many-to-many (M:N) relationship because:

- **One student** can enroll in **many classes**.
- **One class** can have **many students**.

8. Translating Business Rules into Data Model Components

8E. Example: Library System

Consider a library system with the business rule:

- "A book can be borrowed by many members, and a member can borrow many books."

Here:

- **Book** and **Member** are entities.
- The relationship is "borrow."

To determine the relationship type:

- **How many instances of Member are related to one instance of Book?**
 - **Answer:** Many members can borrow the same book (many members to one book).
- **How many instances of Book are related to one instance of Member?**
 - **Answer:** A member can borrow many books (many books to one member).

This is a many-to-many (M:N) relationship because:

- **One book** can be borrowed by **many members**.
- **One member** can borrow **many books**.

Summary

- **Entities** are identified by nouns in business rules.
- **Relationships** are identified by verbs connecting those nouns.
- Relationships are bidirectional, and you determine their type by asking how many instances of one entity are related to one instance of the other.
- Relationships can be one-to-one (1:1), one-to-many (1:M), or many-to-many (M:N), depending on the answers to these questions.

Understanding these relationships is key to constructing an accurate and functional data model, ensuring that the database will effectively support the business's needs.

The Evolution of Data Models

The evolution of data models reflects the ongoing quest to manage data more effectively. Each generation of data models addressed the limitations of previous models while adapting to emerging technological needs and increasing data complexity. Below, I'll explain each major generation of data models in detail, along with further examples to help clarify their significance.

1. Evolution of Major Data Models

1A. First Generation: File System (1960s–1970s)

Example:

- VMS/VSAM

Description:

The first generation of data models was centered around file systems. In these systems, data was stored in flat files, typically in a sequential manner. These files were managed by file management systems such as VMS (Virtual Memory System) and VSAM (Virtual Storage Access Method), commonly used on IBM mainframe systems.

Characteristics:

- **Record Management:** File systems managed data at the record level but lacked the ability to manage relationships between different data entities.
- **Manual Relationships:** Any relationships between data had to be manually implemented by the programmer, usually through additional code or by creating complex file structures.

Shortcomings:

- **Lack of Flexibility:** Since relationships weren't natively supported, making changes to the structure or queries was cumbersome and error-prone.
- **Data Redundancy:** Multiple copies of the same data were often stored in different files, leading to redundancy and inconsistency.

Analogy:

Think of the file system model like a library with a basic card catalog. Each book (record) is listed in the catalog, but if you want to find books on related topics (relationships), you have to manually look them up without any direct connection between the cards.

1B. Second Generation: Hierarchical and Network Models (1970s)

Example:

- IMS, ADABAS, IDS-II

Description:

The second generation introduced the hierarchical and network models, which allowed data to be organized in more complex structures. These models were the first to attempt managing relationships within the database itself, using navigational access methods.

- **Hierarchical Model:** Data is organized in a tree-like structure where each record has a single parent, but potentially many children. IBM's IMS (Information Management System) is a classic example.
- **Network Model:** Data is organized in a graph structure, allowing for more complex relationships where records can have multiple parents and children. ADABAS (Adaptable Database System) and IDS-II (Integrated Data Store) are examples.

Characteristics:

- **Navigational Access:** Data was accessed through predefined paths, making it necessary to know the path to the data before retrieving it.
- **Efficiency:** These models were more efficient than file systems for certain types of queries, especially in applications with well-defined relationships, such as banking or airline reservation systems.

Shortcomings:

- **Complexity:** The navigational approach made it difficult to perform ad-hoc queries or modify the database structure, as the relationships were rigid and predefined.
- **Maintenance Challenges:** The complexity of managing these relationships often led to maintenance challenges as the database grew.

Analogy: Imagine a family tree (hierarchical model) where every person (record) can trace their lineage back to a single ancestor (root). In a social network (network model), however, each person can be connected to multiple other people in various relationships, creating a web of connections.

1C. Third Generation: Relational Model (Mid-1970s)

Example:

- DB2, Oracle, MS SQL Server, MySQL

Description:

The third generation introduced the relational model, which revolutionized data management by organizing data into tables (relations) based on a schema. This model, proposed by E.F. Codd, simplified data management and querying, making databases more accessible to users.

Characteristics:

- **Conceptual Simplicity:** Data is stored in tables with rows (records) and columns (fields), making it easy to understand and work with.
- **Entity Relationship (ER) Modeling:** ER models were used to design the database schema, allowing for a clear representation of entities and their relationships.
- **SQL (Structured Query Language):** SQL became the standard language for querying relational databases, providing powerful and flexible tools for data manipulation.

Shortcomings:

- **Performance Issues:** As databases grew in size and complexity, performance could suffer, particularly with complex joins or large datasets.
- **Limited Data Types:** The relational model primarily supported structured data, making it less suitable for unstructured or semi-structured data like multimedia files or XML.

Analogy: The relational model is like an Excel spreadsheet where each sheet is a table, and you can easily link data from different sheets using common columns (keys). Queries are like formulas that let you pull and manipulate data from different sheets based on your needs.

1D. Fourth Generation: Object-Oriented and Object/Relational Models (Mid-1980s)

Example:

- Versant, Objectivity/DB, DB2 UDB, Oracle

Description:

The fourth generation saw the emergence of object-oriented databases, which integrated object-oriented programming concepts with database management. These databases stored data as objects, similar to how data is managed in object-oriented programming languages.

- **Object-Oriented Model:** Data is stored as objects, which include both data (attributes) and behavior (methods). This model supports complex data types, making it suitable for applications like CAD/CAM, multimedia, and web databases.
- **Object/Relational (O/R) Model:** The O/R model extended relational databases to support object-oriented features, allowing for a mix of relational and object data types.

Characteristics:

- **Rich Data Types:** The object-oriented model could handle complex data types such as images, audio, and video.
- **Star Schema Support:** The O/R model also supported data warehousing structures like star schemas, which optimized complex queries on large datasets.
- **Web Integration:** As the web became more prevalent, these models supported web databases, enabling dynamic content generation and e-commerce applications.

Shortcomings:

- **Complexity:** Object-oriented databases introduced complexity in both design and maintenance, requiring specialized knowledge to manage effectively.
- **Performance Overhead:** The additional complexity sometimes led to performance overhead, especially in applications requiring high-speed transactions.

Analogy: The object-oriented model is like a toolbox where each tool (object) is not just a static item but includes instructions on how to use it (methods). The O/R model is like having a hybrid toolbox where some tools are simple and basic (relational) while others are more complex and multifunctional (object-oriented).

1E. Fifth Generation: XML and Hybrid DBMS (Mid-1990s)

Example:

- dbXML, Tamino, DB2 UDB, Oracle, MS SQL Server, PostgreSQL

Description:

In the mid-1990s, the explosion of the internet and the need to manage unstructured data led to the development of XML databases and hybrid DBMS. These systems combined traditional relational databases with support for XML data, enabling the management of both structured and unstructured data.

Characteristics:

- **XML Support:** These databases natively supported XML documents, allowing for the storage and querying of semi-structured data, which is common in web applications.
- **Hybrid DBMS:** These systems combined relational and object-oriented features with XML support, providing flexibility in data management.
- **Large Database Support:** As data sizes grew into the terabyte range, these systems were designed to handle massive amounts of data efficiently.

Shortcomings:

- **Complexity in Querying:** Querying XML data could be complex, requiring specialized query languages like XPath or XQuery in addition to SQL.
- **Performance Issues:** Managing both structured and unstructured data in a single system could lead to performance challenges, especially with large datasets.

Analogy: Imagine a filing cabinet (relational database) that now has folders with detailed documents (XML) inside. You can still organize the folders by their labels (structured data), but you also need to look inside the folders to understand the full content (unstructured data).

1F. Emerging Models: NoSQL (Early 2000s to Present)

Example:

- SimpleDB (Amazon), BigTable (Google), Cassandra (Apache), MongoDB, Riak

Description:

The emergence of NoSQL databases represents a response to the limitations of traditional relational databases in handling the vast amounts of unstructured and semi-structured data generated by modern applications. NoSQL databases are designed to be distributed, highly scalable, and capable of handling large-scale, high-performance workloads.

Shortcomings:

- **Lack of ACID Compliance:** Many NoSQL databases trade off strict consistency for availability and partition tolerance, which can lead to eventual consistency rather than immediate consistency.
- **Proprietary APIs:** NoSQL databases often come with proprietary APIs, making it difficult to migrate data between different systems or integrate with existing relational databases.

Characteristics:

- **Key-Value Stores:** Simple and fast, these databases store data as key-value pairs, making them ideal for caching and session management.
- **Column Stores:** Optimized for read-heavy workloads, these databases store data by columns rather than rows, which is useful for analytics and big data applications.
- **Document Stores:** Like MongoDB, these databases store data in JSON-like documents, allowing for flexible schemas and easy integration with web applications.
- **Distributed Systems:** NoSQL databases are designed to run on distributed clusters, ensuring high availability and fault tolerance.
- **Scalability:** These databases can handle petabytes of data across thousands of servers, making them suitable for large-scale applications like social networks and e-commerce platforms.

Analogy: NoSQL databases are like different types of warehouses designed for specific purposes. Some are optimized for quick access to individual items (key-value stores), while others are designed to store massive amounts of data in a way that's easy to analyze (column stores). Each type of warehouse has its own strengths, depending on the needs of the application.

1B. Second Generation: Hierarchical and Network Models (1970s)

Example:

- IMS, ADABAS, IDS-II

Description:

The second generation introduced the hierarchical and network models, which allowed data to be organized in more complex structures. These models were the first to attempt managing relationships within the database itself, using navigational access methods.

- **Hierarchical Model:** Data is organized in a tree-like structure where each record has a single parent, but potentially many children. IBM's IMS (Information Management System) is a classic example.
- **Network Model:** Data is organized in a graph structure, allowing for more complex relationships where records can have multiple parents and children. ADABAS (Adaptable Database System) and IDS-II (Integrated Data Store) are examples.

Characteristics:

- **Navigational Access:** Data was accessed through predefined paths, making it necessary to know the path to the data before retrieving it.
- **Efficiency:** These models were more efficient than file systems for certain types of queries, especially in applications with well-defined relationships, such as banking or airline reservation systems.

Shortcomings:

- **Complexity:** The navigational approach made it difficult to perform ad-hoc queries or modify the database structure, as the relationships were rigid and predefined.
- **Maintenance Challenges:** The complexity of managing these relationships often led to maintenance challenges as the database grew.

Analogy: Imagine a family tree (hierarchical model) where every person (record) can trace their lineage back to a single ancestor (root). In a social network (network model), however, each person can be connected to multiple other people in various relationships, creating a web of connections.

2. Hierarchical Model

The **hierarchical model** was one of the earliest database models, designed to manage large volumes of data in complex systems like manufacturing projects. It represents data in a tree-like structure, where each record is called a **segment**.

2A. Tree Structure:

Imagine the hierarchical model as an upside-down tree. The top of the tree is the **root segment**, which can be thought of as the "parent" record. Beneath this root, there are multiple levels of segments that represent data. Each segment in the level below is the "child" of the segment directly above it.

This creates a **1 (one-to-many)** relationship, where one parent segment can have multiple child segments, but each child has only one parent.

2B. Example:

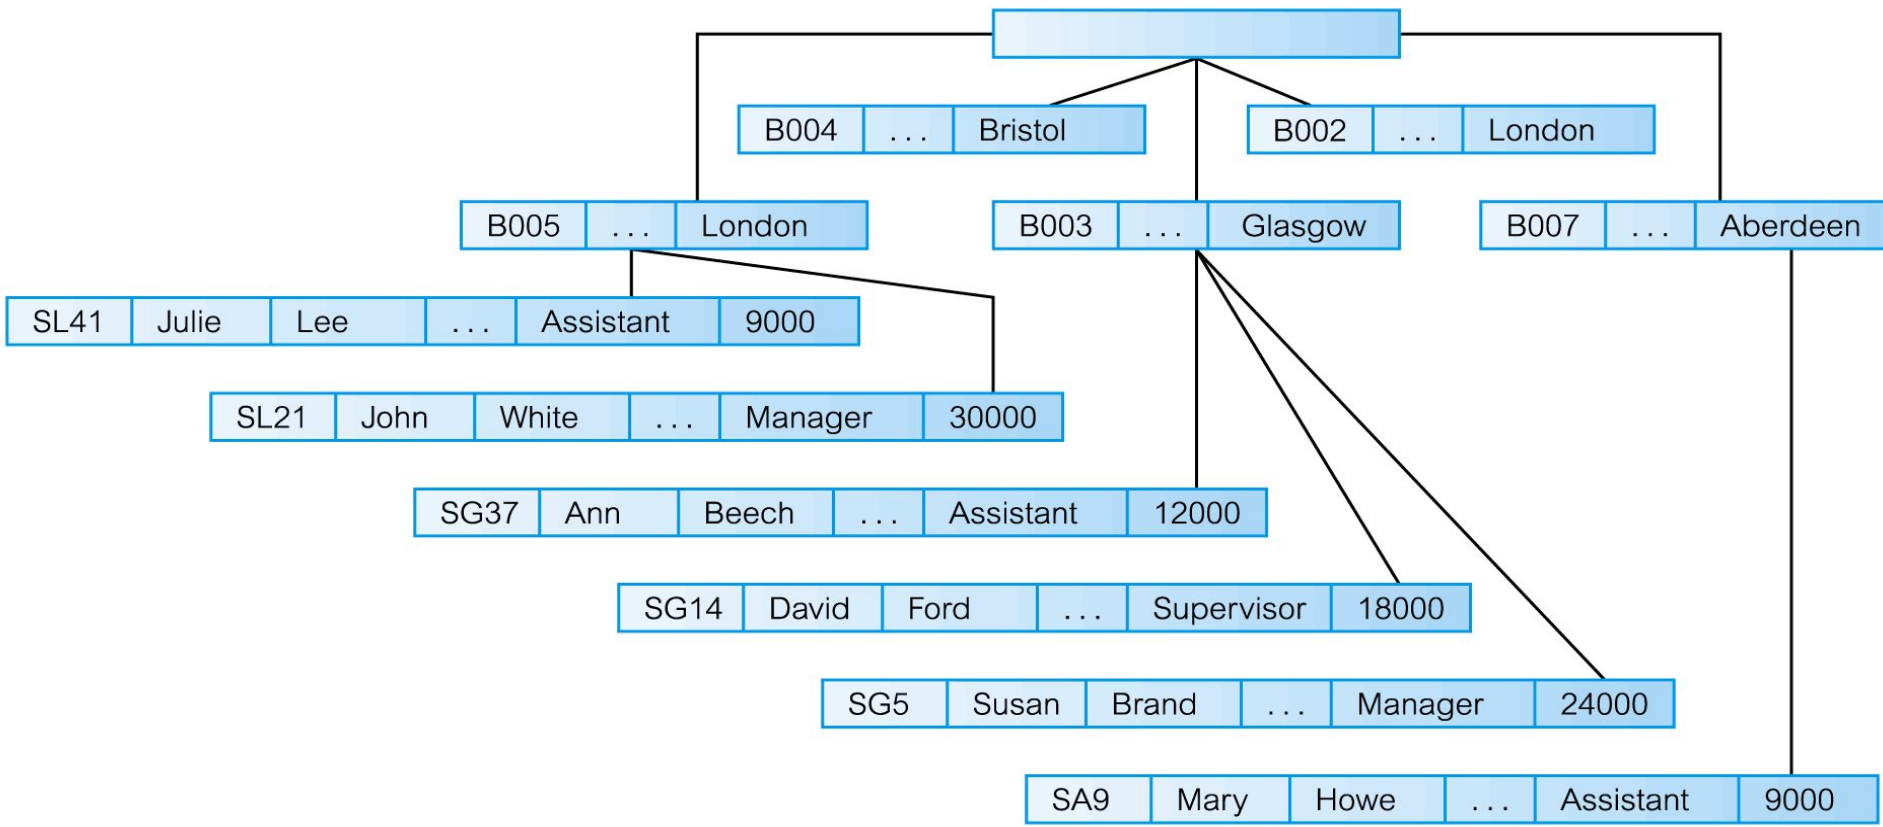
Consider a company's organizational chart as an example. The company's CEO is at the top (root segment). Beneath the CEO, there are several department heads (child segments), and under each department head, there are multiple employees (further child segments). Each department head reports only to the CEO, and each employee reports only to one department head.

2C. Segment:

In this context, each position in the company hierarchy (CEO, department head, employee) represents a segment, similar to a record in a file system.

However, the hierarchical model has limitations. For example, it's difficult to represent relationships that don't fit into a strict hierarchy. If an employee works in multiple departments, the model struggles to represent this since a child segment (employee) can only have one parent segment (department head).

2. Hierarchical Model



3. Network Model

The **network model** was developed to address some of the limitations of the hierarchical model, especially in handling more complex relationships. It allows for more flexible data representation by enabling records to have multiple parent segments.

3A. More Flexible Relationships:

Unlike the hierarchical model, where each child segment can only have one parent, the network model allows a child segment to have multiple parents. This creates a more interconnected web of data, which is why it's called a network model.

3B. Example:

Consider a university database where students are enrolled in multiple courses, and each course has multiple students. In this scenario, both "students" and "courses" are segments. A student can be associated with multiple courses, and each course can have multiple students enrolled. The network model can represent these many-to-many relationships efficiently.

3C. Schema and Subschema:

- **Schema:** In the network model, the **schema** is the overall structure of the database as defined by the database administrator. It defines how all the data is connected and organized.
- **Subschema:** The **subschema** is a portion of the database that an application program can interact with. It's a specific view of the data tailored to a particular application's needs.

3D. Data Manipulation and Definition Languages (DML & DDL):

- **DML (Data Manipulation Language):** This is a set of commands that allows users to interact with the data in the database. Commands like **SELECT**, **INSERT**, **UPDATE**, and **DELETE** are examples of DML. For instance, a user might use a **SELECT** command to retrieve all students enrolled in a particular course.
- **DDL (Data Definition Language):** DDL is used by database administrators to define the structure of the database, including the schema and subschema. For example, an administrator might use a DDL command to create a new table in the database to represent a new type of data, like a "professors" segment.

3. Network Model

Branch Data

B005	22 Deer Rd	London
B007	16 Argyl St	Aberdeen
B003	163 Main St	Glasgow
B004	32 Manse Rd	Bristol
B002	56 Clover Dr	London

Staff Data

SL41	Jul e	Lee	...	Assistant	9000
SL21	John	White	...	Manager	30000
SA9	Mary	Howe	...	Assistant	9000
SG37	Ann	Beech	...	Assistant	12000
SG14	David	Ford	...	Supervisor	18000
SG5	Susan	Brand	...	Manager	24000



3. Network Model

3E. Comparison and Evolution

While both the hierarchical and network models were significant in the early days of database development, they also had drawbacks:

- **Hierarchical Model:** Its strict parent-child relationships made it difficult to represent more complex, interconnected data.
- **Network Model:** Although more flexible than the hierarchical model, it became cumbersome as databases grew in size and complexity. The network model's lack of ad-hoc query capabilities meant that even simple tasks required complex programming, leading to inefficiencies.

Due to these limitations, both models were eventually replaced by the **relational model** in the 1980s, which offered more flexibility, easier querying, and better data independence. The relational model introduced a more straightforward way to handle complex relationships, making it more suitable for the growing data needs of modern applications.

Hint:

As you can see every new model overcome the previous disadvantage of the model and it is to manage.

4. The Relational Model

The **relational model** revolutionized the way data was managed and interacted with, marking a significant departure from the earlier database models that were more cumbersome and complex to use.

Introduced by **E. F. Codd** in 1970, the relational model is grounded in mathematical set theory and introduces a way of structuring and manipulating data that is both powerful and **easy to understand**.

4A. Key Concepts of the Relational Model:

Relation (Table):

- In the relational model, **data is organized into relations**, which are typically visualized as tables. Each table consists of rows and columns.
- **Rows** (also known as **tuples**) represent individual records or entities within the table.
- **Columns** represent attributes or fields that describe the **characteristics of the entities**.
- For example, consider a table named **Employees** with columns like **EmployeeID**, **Name**, **Position**, and **Salary**. Each row in this table would represent a specific employee, with details filled into each corresponding column.

Tuple (Row):

- A tuple is essentially a single row in a table. Each tuple contains data corresponding to each attribute of the table.
- In the **Employees** table example, a tuple might look like **(101, 'John Doe', 'Manager', 75000)**, representing the details of one specific employee.

Attributes (Columns):

- Attributes define the properties or characteristics of the data that is being stored.
- Each column in a table is an attribute, and it holds specific types of data, like integers, strings, dates, etc.

Relations Between Tables:

- **Tables in a relational database can be related to one another through shared attributes. These relationships enable the joining of tables, which allows for complex queries and data retrieval operations.**
- For instance, consider another table named **Departments**, with columns **DepartmentID** and **DepartmentName**. The **Employees** table might have a **DepartmentID** column that corresponds to a **DepartmentID** in the **Departments** table, creating a relationship between these two tables.

4B. Example of the Relational Model in Action:

Imagine you're managing data for a small library. You might have the following tables:

Books table:

- Columns: **BookID**, **Title**, **Author**, **Genre**, **PublishedYear**
- Rows (Tuples): Each row represents a specific book.
 - Example Row: (1, 'To Kill a Mockingbird', 'Harper Lee', 'Fiction', 1960)

Members table:

- Columns: **MemberID**, **Name**, **Address**, **PhoneNumber**
- Rows (Tuples): Each row represents a library member.
 - Example Row: (101, 'Jane Smith', '123 Main St', '555-1234')

Loans table:

- Columns: **LoanID**, **BookID**, **MemberID**, **LoanDate**, **ReturnDate**
- Rows (Tuples): Each row represents a loan of a book to a member.
 - Example Row: (1001, 1, 101, '2024-08-01', '2024-08-15')

In this example:

- The **Books** table holds information about each book in the library.
- The **Members** table holds information about each member of the library.
- The **Loans** table records which books are on loan to which members. It relates the **Books** table and the **Members** table through the **BookID** and **MemberID** columns, respectively.

4C. Advantages of the Relational Model:

Simplicity:

- The relational model's structure is straightforward. Users can interact with data in an intuitive way, without needing to understand the underlying complexities.
- This is akin to the "automatic transmission" analogy mentioned earlier—users don't need to know how the transmission works to drive the car.

Flexibility:

- Because of its basis in set theory, the relational model allows for flexible querying and data manipulation. Users can easily retrieve and update data across multiple tables.

Data Integrity:

- The model supports data integrity through constraints like primary keys (which uniquely identify rows in a table) and foreign keys (which establish and enforce relationships between tables).

Scalability:

- As computing power has increased, relational databases have scaled to handle vast amounts of data, making them suitable for everything from small personal databases to massive enterprise systems.

4D. Example of Querying a Relational Database:

- Suppose you want to find out which members have borrowed "To Kill a Mockingbird." You would perform a query that joins the **Books**, **Loans**, and **Members** tables:

```
SELECT Members.Name, Loans.LoanDate
FROM Members
JOIN Loans ON Members.MemberID = Loans.MemberID
JOIN Books ON Loans.BookID = Books.BookID
WHERE Books.Title = 'To Kill a Mockingbird';
```

- This query fetches the names of members and the dates they borrowed "To Kill a Mockingbird," by joining the relevant tables based on shared attributes (keys).

Conclusion

The relational model's development marked a watershed moment in the history of database management. Its focus on simplicity and its grounding in mathematical theory provided a robust framework that continues to underpin most modern database systems. Today, relational databases are ubiquitous, with RDBMS software like Oracle, SQL Server, and MySQL being used across a variety of industries to manage and query large datasets efficiently and effectively.

4E. Understanding Table Relationships Through a Common Attribute:

In relational databases, tables can be linked or related to one another through a common attribute, often called a foreign key. This allows for data to be organized across multiple tables while still being able to be queried together in meaningful ways. The idea is to minimize redundancy and improve data organization.

Table name: AGENT (first six attributes)

Database name: Ch02_InsureCo

AGENT_CODE	AGENT_LNAME	AGENT_FNAME	AGENT_INITIAL	AGENT_AREACODE	AGENT_PHONE
501	Alby	Alex	B	713	228-1249
502	Hahn	Leah	F	615	862-1244
503	Okon	John	T	615	123-5589

Link through AGENT_CODE

Table name: CUSTOMER

CUS_CODE	CUS_LNAME	CUS_FNAME	CUS_INITIAL	CUS_AREACODE	CUS_PHONE	CUS_INSURE_TYPE	CUS_INSURE_AMT	CUS_RENEW_DATE	AGENT_CODE
10010	Ramas	Alfred	A	615	844-2573	T1	100.00	05-Apr-2022	502
10011	Dunne	Leona	K	713	894-1238	T1	250.00	16-Jun-2022	501
10012	Smith	Kathy	W	615	894-2285	S2	150.00	29-Jan-2023	502
10013	Olowski	Paul	F	615	894-2180	S1	300.00	14-Oct-2022	502
10014	Orlando	Myron		615	222-1672	T1	100.00	28-Dec-2023	501
10015	O'Brian	Amy	B	713	442-3381	T2	850.00	22-Sep-2022	503
10016	Brown	James	G	615	297-1228	S1	120.00	25-Mar-2023	502
10017	Williams	George		615	290-2556	S1	250.00	17-Jul-2022	503
10018	Farriss	Anne	G	713	382-7185	T2	100.00	03-Dec-2022	501
10019	Smith	Olette	K	615	297-3809	S2	500.00	14-Mar-2023	503

Example from the Image:

Table: AGENT

- This table contains information about sales agents, with attributes like AGENT_CODE, AGENT_LNAME (last name), AGENT_FNAME (first name), AGENT_INITIAL, AGENT_AREACODE, and AGENT_PHONE.
- Each agent is uniquely identified by the AGENT_CODE, which acts as a primary key in this table.

Table: CUSTOMER

- This table contains information about customers, with attributes like CUS_CODE, CUS_LNAME (last name), CUS_FNAME (first name), CUS_INITIAL, CUS_AREACODE, CUS_PHONE, CUS_INSURE_TYPE, CUS_INSURE_AMT, CUS_RENEW_DATE, and AGENT_CODE.
- Each customer is uniquely identified by CUS_CODE, which acts as the primary key in this table.
- The AGENT_CODE in the CUSTOMER table is a foreign key that references the AGENT_CODE in the AGENT table, thereby linking the customer to their respective sales agent.

4E. Understanding Table Relationships Through a Common Attribute:

Linking Through a Common Attribute

The **AGENT_CODE** is the common attribute that links the **CUSTOMER** table with the **AGENT** table. Here's how this works:

Customer-Agent Relationship:

- The **AGENT_CODE** in the **CUSTOMER** table points to the **AGENT_CODE** in the **AGENT** table. This linkage allows the database to associate each customer with a specific sales agent.
- For example, **Leona Dunne** has an **AGENT_CODE** of **501**. By looking at the **AGENT** table, we can see that **501** corresponds to **Alex Alby**. Thus, we can determine that **Leona Dunne's** sales agent is **Alex Alby**.

Example Query to Demonstrate Linking:

- Imagine you want to find out which customers are associated with the agent named **Leah Hahn**. You would first find **Leah Hahn's AGENT_CODE** in the **AGENT** table, and then use that code to find all corresponding customers in the **CUSTOMER** table.
- Here's how you could do that with an SQL query:
 - This query joins the **CUSTOMER** and **AGENT** tables on the **AGENT_CODE** attribute.
 - It then filters to show only the customers associated with the agent **Leah Hahn**.

```
SELECT CUSTOMER.CUS_LNAME, CUSTOMER.CUS_FNAME
FROM CUSTOMER
JOIN AGENT ON CUSTOMER.AGENT_CODE = AGENT.AGENT_CODE
WHERE AGENT.AGENT_LNAME = 'Hahn' AND AGENT.AGENT_FNAME = 'Leah';
```


4E. Understanding Table Relationships Through a Common

Attribute:

Minimizing Redundancy:

The relational model is designed to reduce redundancy:

- **Without Relations:** If all data was stored in a single table, the agent information would have to be repeated for every customer they handle, leading to data redundancy.
- **With Relations:** By splitting the data into related tables, the agent information is stored once in the **AGENT** table and referenced by the **CUSTOMER** table as needed. This not only saves space but also makes the data easier to manage. For example, if an agent's phone number changes, it only needs to be updated in one place.

Further Example:

Consider a scenario with a **Products** table and an **Orders** table:

- **Products Table:**
 - Attributes: **ProductID**, **ProductName**, **Price**
- **Orders Table:**
 - Attributes: **OrderID**, **OrderDate**, **CustomerID**, **ProductID**

Here, **ProductID** in the **Orders** table serves as a foreign key linking back to the **Products** table. This setup allows you to store each product's details only once in the **Products** table but reference them multiple times in the **Orders** table, enabling a powerful and efficient way to manage large datasets.

Conclusion:

The relational model's use of common attributes to link tables provides a flexible and efficient way to organize data, enabling complex queries and data retrieval without unnecessary redundancy. By maintaining relationships between tables through keys like **AGENT_CODE**, databases can store data in a highly structured manner, which is both easy to manage and scalable.

4F. Concepts in Relational Diagrams:

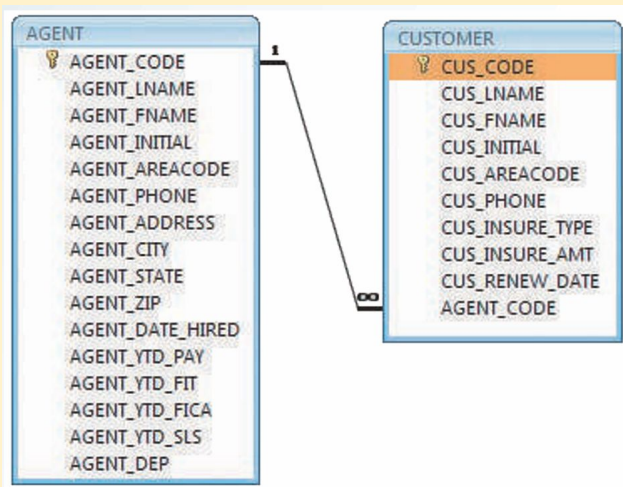


Figure provides a visual representation of a **relational diagram**, which is a key tool in understanding how entities in a relational database are connected. Let's break down the key concepts illustrated by this diagram.

Entities:

- In the diagram, **entities** are represented by the boxes labeled **AGENT** and **CUSTOMER**.
- An **entity** is a real-world object or concept that has data stored about it. In this case, the **AGENT** and **CUSTOMER** are the entities.

Attributes:

- **Attributes** are the pieces of data that describe the entities. For example, **AGENT_CODE**, **AGENT_LNAME**, **AGENT_FNAME**, etc., are attributes of the **AGENT** entity. Similarly, **CUS_CODE**, **CUS_LNAME**, **CUS_FNAME**, etc., are attributes of the **CUSTOMER** entity.
- Attributes are often thought of as the columns in a table.

4F. Concepts in Relational Diagrams:

Primary Key:

- The primary key is the attribute or combination of attributes that uniquely identifies each record in an entity.
- In the **AGENT** table, the primary key is **AGENT_CODE**. In the **CUSTOMER** table, the primary key is **CUS_CODE**. The primary key is usually represented by a key symbol.

Foreign Key:

- A **foreign key** is an attribute in one table that is the primary key in another table. It creates a link between the two tables.
- In the diagram, **AGENT_CODE** in the **CUSTOMER** table is a foreign key that links the **CUSTOMER** entity to the **AGENT** entity.

Relationships:

- The **relationship** between entities defines how they are connected. In a relational diagram, this is often indicated **by a line connecting** the two entities.
- The relationship can be one of the following types:
 - **1:1 (One-to-One)**: One record in an entity is associated with one record in another entity.
 - **1(One-to-Many)**: One record in an entity is associated with many records in another entity. This is the relationship type shown in the diagram.
 - **M(Many-to-Many)**: Many records in one entity are associated with many records in another entity.

4G. Practical Example:

Consider a real-world example in a company:

- **Entity 1: Employee**
 - Attributes might include **EmployeeID** (Primary Key), **LastName**, **FirstName**, **HireDate**, and **DepartmentID**.
- **Entity 2: Department**
 - Attributes might include **DepartmentID** (Primary Key), **DepartmentName**, **ManagerID**.
- **Relationship:**
 - One department (**DepartmentID**) can have many employees (**EmployeeID**), but each employee is assigned to one department.

In a relational diagram, the **DepartmentID** in the **Employee** table would be the foreign key that links back to the **DepartmentID** in the **Department** table, creating a relationship that shows which employees belong to which department.

Conclusion

A **relational diagram** is a powerful tool for visually representing the structure of a relational database. It helps in understanding how entities (tables) are connected through relationships (like one-to-many) and how attributes (columns) are used to link these entities, ensuring that data can be efficiently organized, accessed, and maintained.

4H. Relational Database and Its Components: Detailed Explanation:

A **relational database** is a type of database that stores data in tables, where each table represents a collection of related entities. Think of a table as a matrix of rows and columns:

- **Rows** (also called records) represent individual entries.
- **Columns** (also called fields) represent attributes of the data.

Table vs. File:

A table in a relational database might resemble a traditional file, like a spreadsheet or a CSV (Comma-Separated Values) file, in its appearance and structure, but there's a significant difference:

File:

In a file-based system, the structure and data are often interwoven. If you change the file's structure, you often need to change how you interact with that data.

Table (in a relational database):

The structure (schema) and the data are separate and independent. This means that users or designers are not concerned with how or where the data is physically stored. They focus on the logical design—how data is structured in tables, the relationships between tables, and how queries retrieve this data.

This separation of data from its physical storage is what allows **data independence** and **structural independence**. The user only perceives the logical structure, making it easier to design, manage, and retrieve data.

4H. Relational Database and Its Components: Detailed Explanation:

Rise of the Relational Data Model:

The relational data model became dominant because of its flexibility and the powerful query language it employs—**SQL (Structured Query Language)**.

SQL:

SQL is a declarative language, meaning that it allows users to specify **what** they want to retrieve or manipulate, without having to specify **how** the database should accomplish this task. This is in contrast to procedural programming languages, where the user must write step-by-step instructions to accomplish a task.

For example:

- **SQL Query:** `SELECT * FROM employees WHERE department = 'HR';`
 - This query asks the database to retrieve all columns (*) from the `employees` table where the department is 'HR'.
 - The user doesn't need to know how the database will execute this request (how it will search, sort, or access the data). The SQL engine takes care of this.

Components of an SQL-based Relational Database Application:

End-User Interface:

- This is how the user interacts with the database. The interface might be a form, a webpage, or a software application. When a user inputs a request (e.g., asking for a list of employees in HR), the interface generates the appropriate SQL code behind the scenes.
- Software vendors design these interfaces, but they can also be customized. For example, Microsoft Access allows users to create custom forms that generate SQL queries based on user input.

Collection of Tables:

- All data in a relational database is stored in tables. These tables are logically independent of each other, but they can be related through common fields (known as keys).
- For example, in an e-commerce database:
 - A `Customers` table might store customer details.
 - An `Orders` table might store order information.
 - The `Customers` and `Orders` tables could be related through a `CustomerID` field.

4H. Relational Database and Its Components: Detailed Explanation:

Components of an SQL-based Relational Database Application:

SQL Engine:

- The SQL engine is the part of the DBMS that processes SQL queries. Users generally don't see the SQL engine working, but it is responsible for interpreting and executing all SQL commands, whether they are for retrieving data (SELECT queries), updating data (UPDATE queries), or modifying the structure of the database (ALTER TABLE commands).

Example to Illustrate:

Imagine you have a small business that sells products online. You might use a relational database to manage your data.

- **Tables:**
 - **Products:** Stores details about each product (ProductID, ProductName, Price, etc.).
 - **Customers:** Stores customer information (CustomerID, Name, Email, etc.).
 - **Orders:** Stores each order's details (OrderID, CustomerID, ProductID, OrderDate, etc.).

- **End-User Interface:**

- You might have a simple web form where customers can place orders. When a customer submits the form, the interface generates an SQL query to insert the order into the **Orders** table.

- **SQL Engine:**

- When the order is submitted, the SQL engine processes the query and inserts the data into the relevant tables. If you later want to generate a report showing all orders placed in the last week, you might run an SQL query, like:

```
SELECT * FROM Orders WHERE OrderDate >= '2024-08-19';
```

- The SQL engine would process this request, retrieve the relevant data, and return it to you, without you needing to know how it found and retrieved the data.

5. Entity Relationship Model (ERM)

The **Entity Relationship Model (ERM)** is a framework used in database design to depict relationships between data entities in a conceptual, graphical format.

This model is particularly important because it offers a visual representation of the data structure, making it easier to design, manage, and understand complex databases.

Peter Chen first introduced the ER data model in 1976.

5A. Components of the ER Model

Entities:

- An **entity** represents a real-world object or concept about which data is stored in a database. For example, in a university database, entities might include **STUDENT**, **COURSE**, and **INSTRUCTOR**.
- In an **Entity Relationship Diagram (ERD)**, entities are represented by rectangles. The name of the entity, typically a singular noun, is placed inside the rectangle. For instance, a **STUDENT** entity would be represented by a rectangle labeled "STUDENT".
- **Entity Set**: This refers to a collection of similar entities. For example, all students in a university would constitute a **STUDENT** entity set.
- **Example**:
 - Suppose you have an entity named **BOOK** in a library database. This entity might represent all the individual books in the library. The entity set would be all the **BOOKS** in the library, and each individual book is an instance of the **BOOK** entity.

5A. Components of the ER Model

Attributes:

- **Attributes** are the properties or characteristics of an entity. They provide more information about the entity.
- For instance, the **STUDENT** entity might have attributes like **Student_ID**, **First_Name**, **Last_Name**, and **Date_of_Birth**.
- Attributes are usually represented in an ERD by ovals connected to their respective entities.
- **Example:**
 - For the **BOOK** entity, attributes might include **ISBN**, **Title**, **Author**, **Publication_Year**, and **Genre**. Each attribute provides specific details about a book.

Relationships:

- A **relationship** defines how two or more entities are associated with each other. Relationships can be of different types based on their **connectivity**:
 - **One-to-One (1:1)**: Each instance of entity A is associated with a single instance of entity B, and vice versa.
 - **One-to-Many (1:M)**: An instance of entity A can be associated with multiple instances of entity B, but each instance of entity B is associated with only one instance of entity A.
 - **Many-to-Many (M:N)**: Instances of entity A can be associated with multiple instances of entity B and vice versa.
- **Example:**
 - Consider two entities, **AUTHOR** and **BOOK**. The relationship between them could be described as **One-to-Many**, where one **AUTHOR** can write multiple **BOOKS**, but each **BOOK** is written by one **AUTHOR**.
 - Another example might be a **Many-to-Many** relationship between **STUDENT** and **COURSE**. A student can enroll in many courses, and each course can have many students enrolled.

5B. Entity Relationship Diagram (ERD)

An **Entity Relationship Diagram (ERD)** is a graphical representation of the ER model. It visually represents entities, their attributes, and the relationships among them. The ERD is an essential tool in database design because it provides a clear, structured overview of how data is organized and related within the database.

Example of an ERD

Imagine a simplified ERD for a university database:

Entities:

- **STUDENT**: Attributes might include `Student_ID`, `First_Name`, `Last_Name`.
- **COURSE**: Attributes might include `Course_ID`, `Course_Name`.
- **INSTRUCTOR**: Attributes might include `Instructor_ID`, `Name`.

Relationships:

- **ENROLLS_IN**: A **Many-to-Many** relationship between **STUDENT** and **COURSE**. Each student can enroll in multiple courses, and each course can have multiple students.
- **TEACHES**: A **One-to-Many** relationship between **INSTRUCTOR** and **COURSE**. Each instructor can teach multiple courses, but each course is taught by one instructor.

In the ERD, **STUDENT**, **COURSE**, and **INSTRUCTOR** are depicted as rectangles with their respective attributes connected to them. The relationships (**ENROLLS_IN** and **TEACHES**) are depicted by lines connecting the entities, with the type of relationship (1:1, 1:M, M:N) indicated near the connecting lines.

5C. Importance of the ER Model

The ER model is crucial for several reasons:

- **Conceptual Clarity:** It provides a clear and easily understandable visual representation of the data and its relationships.
- **Foundation for Database Design:** It lays the groundwork for designing relational databases, making it easier to translate conceptual designs into actual database schemas.
- **Efficiency in Database Management:** By clearly defining the relationships and structure of the data, it helps in optimizing queries and database performance.

Conclusion

The **Entity Relationship Model (ERM)** is an indispensable tool in database design, providing a structured and visual approach to understanding and organizing data. Through entities, attributes, and relationships, the ERM helps in creating databases that are both efficient and easy to manage. The use of ERDs further enhances the conceptual clarity, making complex database structures easier to comprehend and implement.

5D. The ER Model Notations

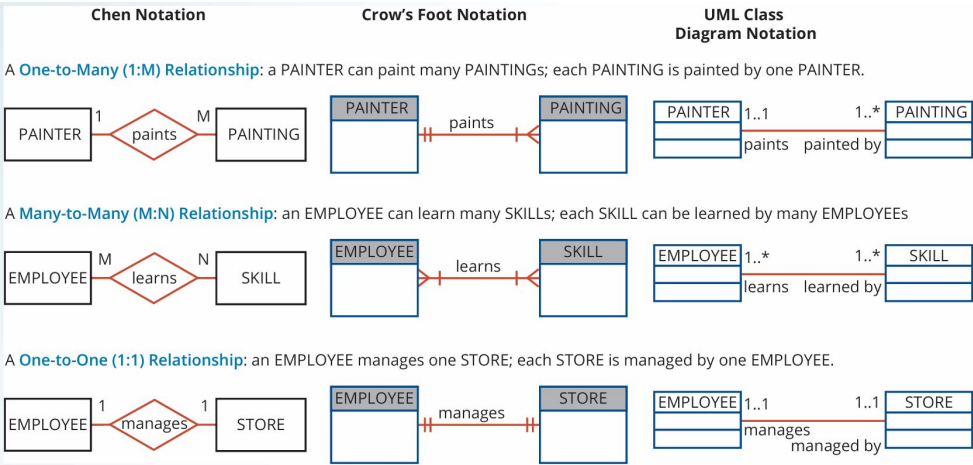


Figure presents different types of relationships in an Entity-Relationship (ER) diagram using three notations: Chen Notation, Crow's Foot Notation, and UML Class Diagram Notation. Here's a detailed explanation of each notation, including examples:

1. Chen Notation:

Explanation:

- This is one of the earliest notations for ER diagrams, introduced by Peter Chen. In Chen notation, entities are represented as rectangles, and relationships are shown as diamonds between the entities. The cardinality (connectivity) is written next to the entities (e.g., 1, M for many, N for many).

Example:

- One-to-Many (1:M) Relationship:** A PAINTER can paint many PAINTINGs, but each PAINTING is painted by one PAINTER.
 - Entities:** PAINTER, PAINTING
 - Relationship:** "paints" (diamond) with a 1 next to PAINTER and M next to PAINTING.
- Many-to-Many (M:N) Relationship:** An EMPLOYEE can learn many SKILLs, and each SKILL can be learned by many EMPLOYEEs.
 - Entities:** EMPLOYEE, SKILL
 - Relationship:** "learns" with an M next to EMPLOYEE and N next to SKILL.
- One-to-One (1:1) Relationship:** An EMPLOYEE manages one STORE, and each STORE is managed by one EMPLOYEE.
 - Entities:** EMPLOYEE, STORE
 - Relationship:** "manages" with a 1 next to both EMPLOYEE and STORE.

5D. The ER Model Notations

2. Crow's Foot Notation:

Explanation:

- Crow's Foot notation is named after the symbol used to represent the "many" side of a relationship, which looks like a crow's foot.
- This notation is highly popular and used in many database design tools. In Crow's Foot notation, the "1" side is represented by a straight line, while the "many" side is represented by a three-pronged "crow's foot."

Example:

- **One-to-Many (1:M) Relationship:** A PAINTER can paint many PAINTINGS.
 - **Representation:** A straight line connects PAINTER to PAINTING on one side, and a crow's foot symbol connects on the many side (PAINTING).
- **Many-to-Many (M:N) Relationship:** An EMPLOYEE can learn many SKILLS.
 - **Representation:** Both EMPLOYEE and SKILL are connected by crow's foot symbols, indicating many-to-many.
- **One-to-One (1:1) Relationship:** An EMPLOYEE manages one STORE.
 - **Representation:** Both EMPLOYEE and STORE are connected by straight lines, indicating a one-to-one relationship.

5D. The ER Model Notations

3. UML Class Diagram Notation:

Explanation:

- UML (Unified Modeling Language) is a versatile diagramming notation used primarily in software engineering. The class diagram notation in UML can also represent ER relationships. Connectivities are indicated by numbers on each end of the relationship line (e.g., 1..*, 1..1).

Example:

- **One-to-Many Relationship:** A PAINTER can paint many PAINTINGS.
 - **Representation:** A line connects PAINTER and PAINTING with "1..1" on the PAINTER side and "1..*" on the PAINTING side.
- **Many-to-Many Relationship:** An EMPLOYEE can learn many SKILLS.
 - **Representation:** A line connects EMPLOYEE and SKILL with "1..*" on both sides
- **One-to-One Relationship:** An EMPLOYEE manages one STORE.
 - **Representation:** A line connects EMPLOYEE and STORE with "1..1" on both sides.

5D. The ER Model Notations

4. Further Example to Illustrate the Differences:

Suppose you have a database for a university where **PROFESSOR** teaches **COURSES**.

- **Chen Notation:**
 - Entities: PROFESSOR, COURSE
 - Relationship: "teaches"
 - Connectivities: 1 next to PROFESSOR, M next to COURSE
- **Crow's Foot Notation:**
 - Similar to the previous example, a straight line would represent PROFESSOR (1 side), and a crow's foot symbol would represent COURSES (many side).
- **UML Class Diagram Notation:**
 - A line connects PROFESSOR and COURSE with "1..1" on the PROFESSOR side and "1..*" on the COURSE side.

This diagram visually explains how different notations can represent the same underlying relationships in a database, but with varying levels of detail and symbols.

Quick Review Question

- What is data model?
- What is the source of business rules?
- What is entity relationship model?

Summary of Main Teaching Points

- Data models are simple representations, usually graphical, of complex real-world data structures.
- Business rules are brief, precise, and unambiguous descriptions of a policies, procedures, or principles within a specific organization.
- Entity relationship models are graphical representation of entities and their relationships in a database structure.